

AI and Deep Learning

# Deep Learning


Jeju National University

Yungcheol Byun

# Agenda

- Merging gates in a computation graph
- Vanishing gradient and ReLU
- MNIST application
- Overfitting and drop-out
- Deep Learning


# Influence of $w$ change on $E$



is multiplication of all the local gradients in the graph  
(chain rule)


# Influence of $w$ change on $E$

Merging gates




# Influence of $w$ change on $E$

Merging gates



# Influence of $w$ change on $E$


Merging gates



is multiplication of **all the local gradients** in the graph  
(chain rule)

# Influence of $w$ change on $E$


Merging gates



is multiplication of **all the local gradients** in the graph  
(chain rule)

# Influence of $w$ change on $E$

Merging gates



Therefore, the local gradient is derivative of the function.



# Influence of $w$ change on $E$


Derivative of  $E$  with respect to  $w$

$$E = (w \cdot 1 - 1)^2$$

$$\frac{\partial E}{\partial w} = \frac{\partial}{\partial w} (w \cdot 1 - 1)^2 = 2(w - 1)$$

# Cost/Error function


for logistic regression



$$hypo = \frac{1}{1 + e^{-wx}}$$

$$E = -y \log(hypo) - (1 - y) \log(1 - hypo)$$

# Computational Graph





$$\frac{\partial E}{\partial w} =$$


# Computational Graph

Merging gates


How we get the local gradient of the merged gate?




# 3-layer NN



# 3-layer NN (simplified)



# Influence of $w$ change on $E$




# 10-layer Neural Network

The giant master, computational graph!



# Influence of $w$ change on $E$



$$\frac{\partial E}{\partial w} = ?$$

*Hint: chain rule!*

# Vanishing Gradient

- The derivative of **sigmoid** function is  $(1 - \text{sigmoid}) * \text{sigmoid}$
- Two multiplication of sigmoid for a single neuron, 20 multiplications for 10 connected neurons
- **Each sigmoid** squashes the input value into the value between 0 and 1.

# Vanishing Gradient

- The influence of  $w$  change on  $E$  is *many* multiplications of the values between 0 and 1, which gives us almost 0.
- Vanishing Gradient
- $w = w - \alpha \cdot$  (almost 0)
- $b = b - \alpha \cdot$  (almost 0)
- Therefore, no change in  $w$  and  $b$

# (Lab) 18.py

- XOR problem using 4-layer neural networks
- Failed owing to vanishing gradient




# The Dark Age in Artificial Intelligence and Neural Networks (~2006)

since back-propagation by Hinton in 1986

# ReLU

proposed by Hahnloser in 2000 and  
demonstrated for deep networks  
in 2011

using **ReLU** (Rectified Linear Unit)  
activation function instead of  
sigmoid



# (Lab) 19.py

- Solving vanishing gradient problem using ReLU activation function
- Back-propagation is working by using ReLU.

So, now can go deeper.



# MNIST







Modified National Institute of  
Standards and Technology  
(USA)



# MINIST





# (Lab) 20.py

- 60,000 training images + 10,000 testing images
- Input image :  $28 * 28$  pixels  $\rightarrow$  784 pixels
- 784 dimension
- 10 classes (output: 0 ~ 9)
- Softmax
- 90.23% of recognition rate


Input Layer


Fully-connected

Output Layer

784


10





# (Lab) 21.py

- Deep Neural Network (4-layer)
- ReLU
- 94.55% accuracy




# (Lab) 22.py

- Applying **initialization** method for  **$w$**  and  **$b$** , not randomly
- **97.23%** of accuracy



# (Lab) 23.py


- Applying initialization method for  $w$  and  $b$ , not randomly
- 6-layer deep neural networks
- 97.83% of accuracy



# Overfitting and drop-out

- The deeper the network is, the more the decision boundary is complex.
- Good at learning data but errors for testing data → **overfitted** to learning data
- Making it less complex by **drop-out** some neurons while learning.

Which do you think is desirable decision boundary?




*While the black line fits the data well,  
the green line is overfit.*


<https://elitedatascience.com>

# Regularization: **Dropout**

“randomly set some neurons to zero in the forward pass”



(a) Standard Neural Net



(b) After applying dropout.

*[Srivastava et al., 2014]*


# (Lab) 24.py

- Applying dropout
- 98.13% of recognition accuracy

# How to Prevent Overfitting


- Train with more data
- Reduce features
- Early stopping
- Ensemble
- regularization

## Early stopping







## Early stopping



Fully-connected

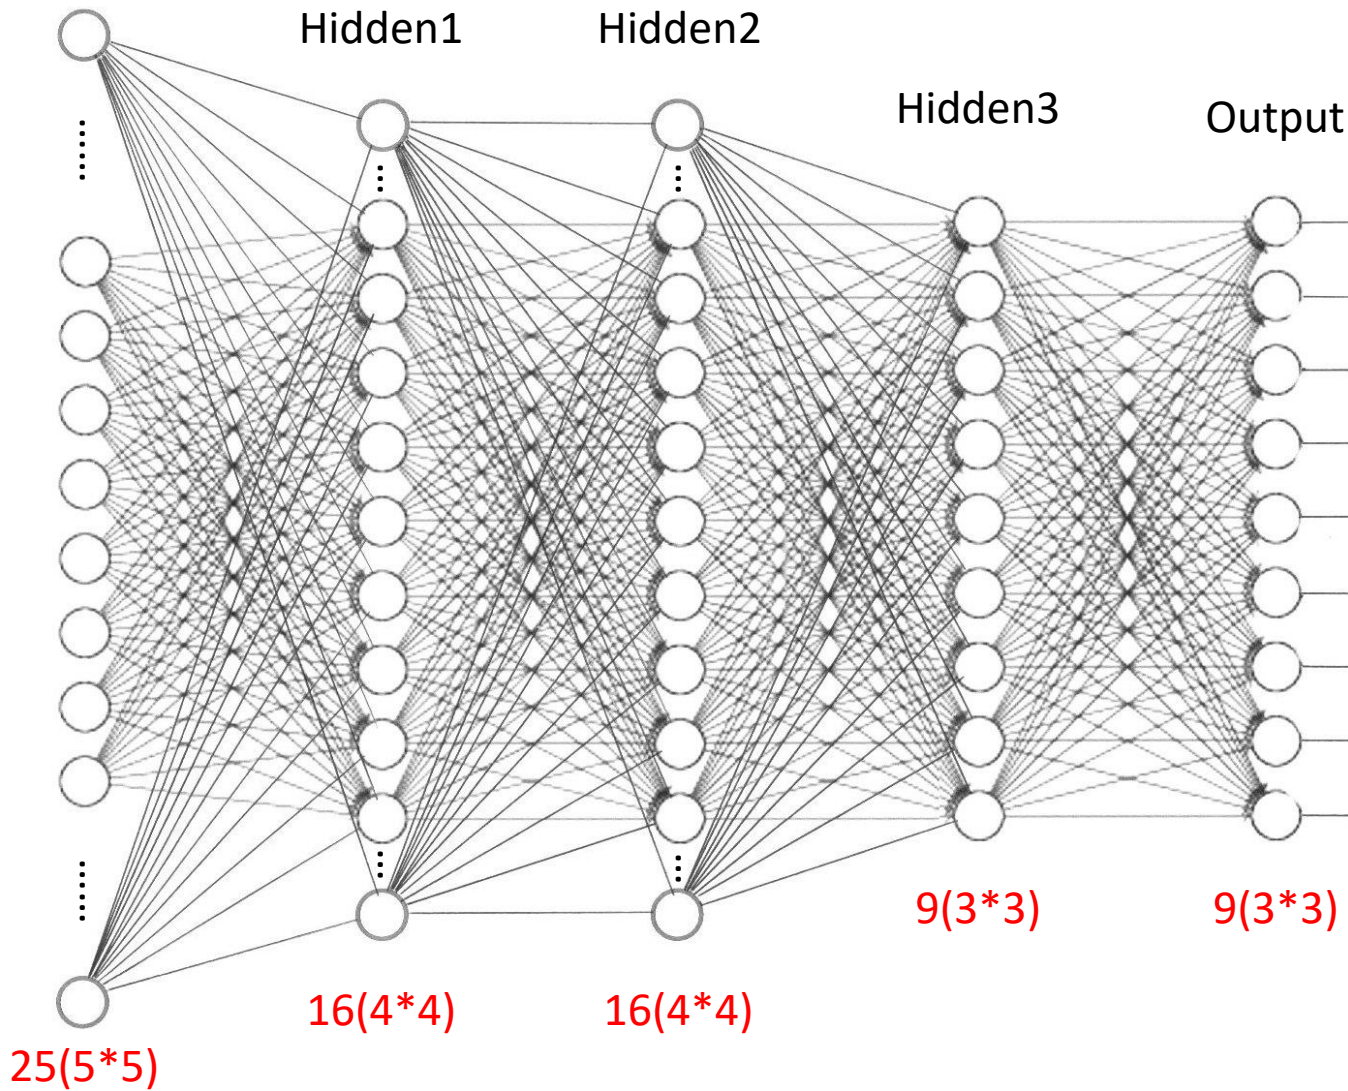


Fully-connected



Input

Fully-connected




Fully connected, then how many synapses(parameters) are there?

$$25 * 16 + 16 * 16 + 16 * 9 + 9 * 9 = 881$$






Fully-connected



Fully connected, so how many connections are there?

$$25 * 16 + 16 * 16 + 16 * 9 + 9 * 9 = 881$$







Geoffrey Hinton, Yann LeCun, Yoshua Bengio, Andrew Ng





# Deep Learning

- in early 2000s (2006, 2010, 2012)
- Deep Neural Networks
- Activation functions (ReLU)
- Weight initialization methods
- Dropout (2014)
- Big data
- GPU

Fully-connected

# FCNN

Any problem?